

# ADR-2: Database and Schema Decision

## 1 CONTEXT

---

The DECODER API requires persistent storage for sensor readings, buildings, and users. A decision must be made on:

1. Database type (SQL vs NoSQL)
2. Specific database technology
3. Schema design
4. Data access strategy (ORM vs raw SQL)

## 2 DECISION

---

I have chosen to use SQLite as the database with JPA/Hibernate for object-relational mapping, using a relational schema with three main tables: readings, buildings, and users.

## 3 RATIONALE

---

### 3.1 WHY SQLITE?

#### 3.1.1 Simplicity for Research Slice

- No separate database server required
- Single file database (decoder.db)
- Perfect for local development and demonstration
- Zero configuration needed

#### 3.1.2 Easy Deployment

- Database file can be included or generated
- No external dependencies for basic setup
- Can easily switch to PostgreSQL for production

#### 3.1.3 SQL Compliance

- Uses standard SQL syntax
- Easy to migrate to PostgreSQL later
- Supports all required features (foreign keys, indexes, etc.)

## 3.2 WHY RELATIONAL SCHEMA?

### 3.2.1 Structured Data

- Sensor readings have clear relationships (building → readings)
- Users have clear relationships (user → buildings)
- Relational model fits the domain naturally

### 3.2.2 Query Requirements

- Need to query readings by building and time range
- Need to filter buildings by owner
- Relational SQL handles these queries efficiently

## 3.3 DATA INTEGRITY:

- Foreign key constraints ensure data consistency
- Referential integrity prevents orphaned records
- Transaction support for atomic operations

## 3.4 WHY JPA/HIBERNATE?

### 3.4.1 Spring Boot Integration

- Seamless integration with Spring Boot
- Automatic transaction management
- Repository pattern support

### 3.4.2 Productivity

- Reduces boilerplate code
- Automatic schema generation (ddl-auto: update)
- Type-safe queries with method names

### 3.4.3 Flexibility

- Can write custom queries when needed
- Easy to switch between databases
- Good for research/prototyping

## 3.5 SCHEMA DESIGN

### 3.5.1 Tables

#### 3.5.1.1 *users*

sql

```
CREATE TABLE users (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    username VARCHAR(255) UNIQUE NOT NULL,
```

```
    role VARCHAR(50) NOT NULL -- ADMIN or OWNER  
);
```

### 3.5.1.2 *buildings*

sql

```
CREATE TABLE buildings (  
    id INTEGER PRIMARY KEY AUTOINCREMENT,  
    name VARCHAR(255) NOT NULL,  
    owner_id INTEGER NOT NULL,  
    address VARCHAR(255),  
    FOREIGN KEY (owner_id) REFERENCES users(id)  
);
```

### 3.5.1.3 *readings*

sql

```
CREATE TABLE readings (  
    id INTEGER PRIMARY KEY AUTOINCREMENT,  
    building_id INTEGER NOT NULL,  
    sensor_id VARCHAR(255) NOT NULL,  
    timestamp TIMESTAMP NOT NULL,  
    value REAL NOT NULL,  
    FOREIGN KEY (building_id) REFERENCES buildings(id)  
);
```

```
CREATE INDEX idx_building_timestamp ON readings(building_id, timestamp);
```

```
CREATE INDEX idx_sensor_timestamp ON readings(sensor_id, timestamp);
```

## 3.6 DESIGN DECISIONS

### 3.6.1 Normalization: Third normal form

- Buildings separate from users (many-to-one)
- Readings separate from buildings (many-to-one)
- Avoids data duplication

### **3.6.2 Indexes**

- Index on `(building\_id, timestamp)` for efficient time-range queries
- Index on `(sensor\_id, timestamp)` for sensor-specific queries

### **3.6.3 Data Types**

- `REAL` for sensor values (allows decimals)
- `TIMESTAMP` for timestamps (ISO 8601 format)
- `INTEGER` for IDs (auto-increment)

### **3.6.4 Constraints**

- Foreign keys ensure referential integrity
- NOT NULL constraints prevent missing data
- UNIQUE on username prevents duplicate users

## **3.7 CONSEQUENCES**

### **3.7.1 Positive**

- Simple Setup: No database server required
- Fast Development: Database auto-creates schema
- Easy Migration: SQLite schema can be migrated to PostgreSQL
- Type Safety: JPA provides compile-time type checking
- Query Efficiency: Indexes support fast time-range queries

### **3.7.2 Negative**

- Concurrency: SQLite has limited concurrent write support
- Scalability: Not suitable for high-volume production use
- No Advanced Features: Missing some PostgreSQL features (full-text search, etc.)

## **3.8 ALTERNATIVES CONSIDERED**

### **3.8.1 PostgreSQL**

Considered but not chosen for initial implementation:

- Requires separate database server
- More complex setup for research slice
- Future: Can be switched by changing datasource URL

### **3.8.2 MongoDB (NoSQL)**

Rejected because:

- Sensor readings have structured schema (better fit for SQL)
- Need complex queries (time ranges, joins)
- JPA/Hibernate don't work well with NoSQL
- Relational model fits the domain better

### **3.8.3 In-Memory Database (H2)**

Considered for tests only:

- Used in integration tests (H2)
- Not suitable for production (data lost on restart)
- SQLite provides persistence while maintaining simplicity

### **3.8.4 Migration Path**

If moving to PostgreSQL:

#### **3.8.4.1 Schema Migration**

yaml

```
spring:  
  datasource:  
    url: jdbc:postgresql://localhost:5432/decoder  
    driver-class-name: org.postgresql.Driver
```

#### **3.8.4.2 Dialect Change**

yaml

```
hibernate:  
  dialect: org.hibernate.dialect.PostgreSQLDialect
```

#### **3.8.4.3 SQL Adjustments**

- Change `INTEGER PRIMARY KEY AUTOINCREMENT` → `SERIAL PRIMARY KEY`
- Adjust timestamp handling if needed
- SQLite-specific syntax may need adjustment

## **3.9 DATA MIGRATION**

- Export SQLite data
- Import into PostgreSQL
- Verify data integrity

## **3.10 TESTING STRATEGY**

Unit Tests: Use in-memory H2 database

- Integration Tests: Use H2 for faster test execution
- Production: SQLite (or PostgreSQL for scale)